

THE GEORGE
WASHINGTON
UNIVERSITY

WASHINGTON, DC

Internet of Things

Lecture 6: Logistic Regression with
PySpark



LOGISTIC REGRESSION





Background

- We want to learn about Logistic Regression as a method for **Classification**.
- Some examples of classification problems:
 - Spam versus “Ham” emails
 - Loan Default (yes/no)
 - Disease Diagnosis
- Above were all examples of Binary Classification





Background


- The convention for binary classification is to have two classes 0 and 1.
- Let's walk through the basic idea for logistic regression.
- We'll also explain why it has the term regression in it, even though it's used for classification!





Background


- Imagine we plotted out some categorical data against one feature.





Background


- The X axis represents a feature value and the Y axis represents the probability of belonging to class 1.





Background


- We can't use a normal linear regression model on binary groups. It won't lead to a good fit:





Background


- We need a function that will fit binary categorical data!





Background


- It would be great if we could find a function with this sort of behavior:





Sigmoid Function


- The Sigmoid (aka Logistic) Function takes in any value and outputs it to be between 0 and 1.





Sigmoid Function


- This means we can take our Linear Regression Solution and place it into the Sigmoid Function.





Sigmoid Function


- This results in a probability from 0 to 1 of belonging in the 1 class.





Sigmoid Function


- We can set a cutoff point at 0.5, anything below it results in class 0, anything above is class 1.





Review

- We use the logistic function to output a value ranging from 0 to 1. Based off of this probability we assign a class.





Model Evaluation

- After you train a logistic regression model on some training data, you will evaluate your model's performance on some test data.
- You can use a confusion matrix to evaluate classification models.





Model Evaluation

- The main point to remember with the confusion matrix and the various calculated metrics is that they are all fundamentally ways of comparing the predicted values versus the true values.
- What constitutes “good” metrics, will really depend on the specific situation!





Model Evaluation

- We can use a confusion matrix to evaluate our model.
- For example, imagine testing for disease.

n=165	Predicted: NO	Predicted: YES
	Actual: NO	Actual: YES
	50	10
	5	100

Example: Test for presence of disease
NO = negative test = False = 0
YES = positive test = True = 1





Confusion Matrix

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Basic Terminology:

- True Positives (TP)
- True Negatives (TN)
- False Positives (FP)
- False Negatives (FN)





Confusion Matrix

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Accuracy:

- Overall, how often is it **correct**?
- $(TP + TN) / \text{total} = 150/165 = 0.91$





Confusion Matrix

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Misclassification Rate
(Error Rate):


- Overall, how often is it **wrong**?
- $(FP + FN) / \text{total} = 15/165 = 0.09$





Confusion Matrix

Type I error
(false positive)



Type II error
(false negative)





Confusion Matrix

		predicted condition	
total population		prediction positive	prediction negative
true condition	condition positive	True Positive (TP)	False Negative (FN) (type II error)
	condition negative	False Positive (FP) (Type I error)	True Negative (TN)





Confusion Matrix

		predicted condition		
		total population		
true condition		prediction positive	prediction negative	Prevalence = $\frac{\Sigma \text{ condition positive}}{\Sigma \text{ total population}}$
	condition positive	True Positive (TP)	False Negative (FN) (type II error)	True Positive Rate (TPR), Sensitivity, Recall, Probability of Detection $= \frac{\Sigma \text{ TP}}{\Sigma \text{ condition positive}}$
	condition negative	False Positive (FP) (Type I error)	True Negative (TN)	False Positive Rate (FPR), Fall-out, Probability of False Alarm $= \frac{\Sigma \text{ FP}}{\Sigma \text{ condition negative}}$
		Positive Predictive Value (PPV), Precision $= \frac{\Sigma \text{ TP}}{\Sigma \text{ prediction positive}}$	False Omission Rate (FOR) $= \frac{\Sigma \text{ FN}}{\Sigma \text{ prediction negative}}$	Positive Likelihood Ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$
		Accuracy $= \frac{\Sigma \text{ TP} + \Sigma \text{ TN}}{\Sigma \text{ total population}}$	False Discovery Rate (FDR) $= \frac{\Sigma \text{ FP}}{\Sigma \text{ prediction positive}}$	Negative Predictive Value (NPV) $= \frac{\Sigma \text{ TN}}{\Sigma \text{ prediction negative}}$
				Negative Likelihood Ratio (LR-) $= \frac{\text{FNR}}{\text{TNR}}$

